


# ECE69500: INFERENCE & LEARNING IN GENERATIVE MODELS

Spring 2021

## Homework 4

The attached CSV file, `HMMdata.csv`, contains three columns of data: one column of class identities ( $\mathbf{X}$ ), and two columns of two-dimensional “features” ( $\mathbf{Y}$ ).

1. Plot the data, with each class in a different color, like this:



2. Assume the data  $\mathbf{Y}$  were generated by a **Gaussian mixture model** with three classes and are observed, while the class identities  $\mathbf{X}$  are not. In whatever programming language you prefer, implement (from scratch) inference under the GMM,

$$q(\mathbf{X}) = \text{Categ}[\boldsymbol{\pi}]$$
$$q(\mathbf{Y}|\mathbf{X}) = \mathcal{N}(\mathbf{C}\mathbf{X}, \boldsymbol{\Sigma})$$

with the parameter values

$$\boldsymbol{\pi} = \left( \frac{3}{16}, \frac{5}{16}, \frac{8}{16} \right)^T, \quad \mathbf{C} = \begin{pmatrix} -2 & 2 & -2 \\ 0 & 2 & 2 \end{pmatrix}, \quad \boldsymbol{\Sigma} = \begin{pmatrix} 1 & -\frac{1}{5} \\ -\frac{1}{5} & 2 \end{pmatrix}$$

Note that we are assuming a fixed covariance matrix across all classes; and that we are treating  $\mathbf{X}$  as a one-hot vector, so that it selects a column of  $\mathbf{C}$  for the mean of the emission.

Thus your code should return a set of posterior probabilities over the latent state  $\mathbf{X}$ :  $q(\mathbf{X}|\mathbf{y})$  for all observations  $\mathbf{y}$ . (As a sanity check, selecting the most probable state will yield an average accuracy of 80%, i.e. `np.mean(np.argmax(qGMM, 0) == X)` is 0.8.)

3. Now assume the data  $\mathbf{Y}$  were generated by a **hidden Markov model** with three classes and are observed, while again the class identities  $\mathbf{X}$  are not. In whatever programming language you prefer, implement (from scratch) inference under the HMM,

$$\begin{aligned} q(\mathbf{X}_1) &= \text{Categ}[\boldsymbol{\pi}] \\ q(\mathbf{X}_n | \mathbf{X}_{n-1}) &= \text{Categ}[\mathbf{A}\mathbf{X}_{n-1}] \\ q(\mathbf{Y}_n | \mathbf{X}_n) &= \mathcal{N}(\mathbf{C}\mathbf{X}_n, \boldsymbol{\Sigma}) \end{aligned}$$

with the parameter values

$$\boldsymbol{\pi} = \left( \frac{3}{5}, \frac{1}{5}, \frac{1}{5} \right), \quad \mathbf{C} = \begin{pmatrix} -2 & 2 & -2 \\ 0 & 2 & 2 \end{pmatrix}, \quad \mathbf{A} = \begin{pmatrix} \frac{4}{10} & \frac{2}{10} & \frac{1}{10} \\ \frac{3}{10} & \frac{5}{10} & \frac{2}{10} \\ \frac{3}{10} & \frac{3}{10} & \frac{7}{10} \end{pmatrix} \quad \boldsymbol{\Sigma} = \begin{pmatrix} 1 & -\frac{1}{5} \\ -\frac{1}{5} & 2 \end{pmatrix}$$

(Notice that  $\boldsymbol{\pi}$  is different from the GMM.)

Thus your code should return a set of posterior probabilities over the latent state  $\mathbf{X}_n$ . (As a sanity check, selecting the most probable state at each time step from the *filter distribution* will yield an average accuracy of 81.5%; and selecting the most probable state at each time step from the *smoother distribution* will yield an average accuracy of 83%.)